

The **Essence of SQL:**

**A Guide to Learning Most of SQL
in the Least Amount of Time**



||

THE ESSENCE OF SQL:

A Guide to Learning Most of SQL
in the Least Amount of Time

David Rozenshtein

Department of Computer Science
Long Island University

||



40087 Mission Boulevard, Suite 167, Fremont, CA 94539

(800) 943-9300 Voice

(510) 654-5116 Fax

<http://www.sqlforum.com>
Library of Congress Catalog Card Number 96-092799

ISBN 0-9649812-1-1

a:> 1996 SQL Forum Press

Editor: Tom Bondur

Cover concept Anna Bondur

Cover design: Tom McLeith

Inside layout Mimi Fujii



4CX87MissionBoulevard,Suill!167, Pmnmt. CA 94539

(800) 9e9300 Voice

(510) 656-5116 Fu

<http://www.sqlforum.com>

..... THE ESSENCE OF SQL: A

Guide to Learning Most of SQL In the Least Amount of Time

Contents

1	Introduction	1
2	Example Database	3
3	A Bird's Eye View of SQL	5
4	The Standard Questions	9
6	The First Two Questions	11
8	Standard Questions Involving "or" and "both-and"	17
7	Negation in SQL	23
8	Standard Questions Involving "at least"	27
9	Negation Revisited	29
10	JYpe 2 SQL Queries	31
11	Using JYpe 2 to Implement Real Negation	35
12	Posing Questions Involving "at most," "exactly," "only," and "either-or"	37
13	Computing Extremes as JYpe 2 Queries with Negation	43
14	A Look	

at SQL Enhancements	45	15 Handling Composite Keys	53	18 Negation
Over Composite Keys	55	17 Standard Questions Involving "every"	59	18
A Brief Interlude	61	19 Additional Standard Questions	63	20 Computing
Aggregates for Groups	65	21 Combining Scalar Expressions with		
Aggregation	73	22 Some Old Questions Revisited	75	23 Global
Aggregations	77	24 Comparing Base Values to A		
		ggregates	79	



Contents

25 Combining type 2 Queries with Aggregation	83	26 The Overall Order of
Evaluation	87	27 NULL Values in SQL
91	28 Introduction to type 3	
SQL Queries	97	
29 Inter-Query Connectors Revisited:		
Existential Quantifiers in SQL	101	30 Some Old Questions
Revisited-Again	103	31 Data Manipulation Facilities of SQL
109	32	
Extensions to SQL	115	33 Last Remarks
117		

List of Figures

Figure2.1: Example database. 3 Figure3.1: DDL for table Student. 5 Figure3.2: SQL query types. 7 Figure4.1: A list of standard questions. 9 FigureS.1: Query Qt for question E1 - Who takes CS112? 12 FigureS.2: The Type 1 SQL evaluation mechanism. 13 Figure5.3: Query Q2 for question E2 - What are student numbers and names of students who take CS112? 14 Figure5.4: Query Q2 rewritten with all explicit prefixes. 16 Figure6.1: Query Q3 for question E3 - Who takes CS112 or CS114? 17 Figure6.2: First attempt at question E4 - Who takes both CS112 andCS114? 18 Figure6.3: A sketch of the query for E4 - Who takes both CS112 andCS114? 20 Figure6.4: Query Q4 for question E4 - Who takes both CS112

andCS114? 21
 Figure6.5: Q4 rewritten with both copies of Take aliased. 22 Figure7.1: Orthogonality
 of questions ES and E6. 24 Figure7.2: Query Q6 for question E6 - Who takes a course
 which is not CS112? 24
 Figure8.1: Query Q7 for question E7 - Who takes at least 2 courses? 27 Figure8.2: The
 query for question: Who takes at least 3 courses? 27 Figure9.1: Questions E8 through
 E13 rephrased to explicitly show negation. 30
 Figure 10.1: A sample Type 2 SQL query. 31 Figure 10.2: The definition of a Type 2
 query. 33 Figure 10.3: The Type 2 SQL evaluation mechanism. 33 Figure 11.1: A
 mystery 'JYpe 2 SQL query involving negation. 35 Figure 11.2: Using Take in both
 FROM clauses. 36 Figure 12.1: Query Q8 for question E8 - Who takes at most 2
 courses? 37 Figure 12.2: Query Q9 for question E9 - Who takes exactly
 2 courses? 38
 Figure 12.3: Query Q10 for question E10 - Who takes only CS112? 39 Figure 12.4:
 Query of Figure 12.3 simplified. 40 Figure 12.5: Query Q11 for question E11 - Who takes
 either CS112
 orCS114? 40
 Figure 13.1: Query Q12 for question E12 - Who are the youngest
 students? 43
 Figure 14.1: Example of query with the omitted WHERE clause. 45 Figure 14.2: Using
 symbol• to mean "all columns." 46



IV: Use of Figures 14.1-14.6

Figure 14.3: Suppressing duplicates from the answer. 47 Figure 14.4: Sorting
 facilities of SQL 48 Figure 14.5: Example use of arithmetic in SQL. 50 Figure 14.6:
 Formatting an answer. 51 Figure 15.1: Query for the question What are full names and
 ages of professors who teach CS112? 53
 Figure 15.2: Query for the question Which courses are taught by at least 2 professors?
 54
 Figure 16.1: First attempt at the question Who does not teach CS112? 55
 Figure 16.2: Example data for tables Professor and Teach. 56 Figure 16.3: Another
 example data for tables Professor and Teach. 57 Figure 16.4: Desired query form for
 the question Who does not teach CS112? 57
 Figure 16.5: Using string concatenation expressions for the
 question Who does not teach CS112? 58
 Figure 17.1: Query Q13 for question E13- Who takes every course? 59 Figure 19.1: A
 list of additional standard questions motivating SQL aggregation facilities. 63
 Figure 20.1: Query Q14 for question E14- For each department that has more
 than 3 professors older than SO, what
 is the average salary of such professors? 65
 Figure 20.2: The evaluation mechanism for queries involving aggregation. 67
 Figure 20.3: A query showing the use of all aggregate functions, and also
 showing the use of Boolean expressions in
 the HAVING clause. 68
 Figure 20.4: The query for the question For each department, what is the
 average salary of those professors
 who are older than SO? 69

Figure 20.5: SQL query for the question For each department/ rank combination that has more than 3 professors older than SO, what is the average salary of such professors? 70

Figure 21.1: Query QIS for question EIS -What is the grade point average (GPA) of each student? 73

Figure 22.1: A query with the GROUP BY and HAVING clauses for the question Who takes at least 3 courses? 75

Figure 23.1: Query QI6 for question E16 - What is the overall average salary of all professors who are older

than SO? 77

Figure 24.1: Query Q17 for question E17 - Whose

(i.e., which professors') salary is greater than the overall average salary? 79

Figure 24.2: Alternative query form for question E17 - Whose (i.e., which

professors') salary is greater than the

overall average salary? 81



THE ESSENCE OF SQL: A Guide to Learning Most of SQL in the Least Amount of Time

Figure 24.3: Query Q18 for question EIS-Whose salary is greater than the average salary within that professor's

department? 81

Figure 25.1: Alternative query for question E12 - Who are the youngest students? - showing the use of

aggregates inside the subqueries. 83

Figure 25.2: Alternative query for question E17-Whose salary is greater than the overall average salary?- showing

the use of binary comparators in front of subqueries. 84

Figure 26.1: A query involving most of the features presented so far. 87 Figure 27.1: Truth tables for the three-valued logic. 92 Figure 27.2: A query that will retrieve names of all professors or will it? 95

Figure 28.1: A sample 'type 3 query. 97 Figure 28.2: A 'type 3 query with two correlations, for the question Who teaches CS112? 99

Figure 29.1: An example query with EXISTS. 101 Figure 30.1: A 'type 3 query for question E9 Who takes exactly 2courses? 103

Figure 30.2: A 'type 3 query for question E12 - Who are the

youngest students? 104

Figure 30.3: A 'type 3 query for question E18 -Whose salary is greater than the average salary within that

professor's department? 105

Figure 30.4: A 'type 3 query for question E13-Who takes every course? 106

Figure 30.5: The middle query from Figure 30.4 after substitution for S.Sno. 106

Figure 30.6: A 'type 3 query with a 'type 2 subquery for question E13 - Who takes every course? 107

Figure 30.7: The middle query from Figure 30.6 after substitution for S.Sno. 108

Figure 31.1: Inserting individual rows. 109 Figure 31.2: Inserting a query result. 110

Figure 31.3: Example of a DELETE command. 111 Figure 31.4: Example of an UPDATE command. 112 Figure 33.1: Assorted reminders for programming in SQL 118



1 Introduction

This essay is dedicated to the proposition that one can often accomplish 80% of the task in 20% of the time, or, in this case, that one can learn most of the important features and

capabilities of SQL quickly. To achieve this, we have developed a new approach to presenting SQL.

Traditionally, SQL is presented in a "bottom-up" fashion, by enumerating all of its features in some order and illustrating each feature with an example query. While complete in its treatment of the language, this approach often fails to clearly distinguish between what is essential and what is secondary in the language, and frequently leaves an impression of too long a "laundry list" of language features. Even more importantly, by concentrating on the language first and motivation

second, this approach does not clearly explain why certain features are present in the language, when they should be used, nor how to actually pose queries using SQL.

_____ ;  : 11

Our approach, on the other hand, is "top-down." We begin by identifying a rather short list of standard questions, or more precisely types of questions, that are often asked of relational databases. We then show how these standard questions are posed in SQL, introducing and motivating the use of its capabilities and features as they become relevant.

As a result, we develop concrete solutions for a set of well specified question types. Solutions for particular business problems can then be developed by selecting the appropriate question type or combination of types and rephrasing the corresponding standard solutions to fit the particular problem domain.

It is important to note that we do not claim completeness in our approach- some features of the language will not be covered here. We also do not claim objectivity in the selection of features. This essay is based on the SQL working and teaching experience of the author, and reflects his views on the language.

While we assume a familiarity with the basic relational concepts of tables, attributes, keys, etc., no prior knowledge of SQL itself is assumed. In fact, those who already know some SQL should set their knowledge aside and start anew with this essay.

Finally, whatever liberties we take in this essay are all motivated by its primary purpose - to leave the reader with a deep understanding of the essence of SQL and a

2 Example Database Database

All of the SQL queries in this essay will be posed with respect to the example database shown in Figure 2.1. This database consists of five tables, three of which correspond to entity classes: Student, Course and Professor, and two to relationships: Take- between students and courses, and Teach- between professors and courses.

Student(SnQ, Sname, Age)



Course(Cno, Title, Credits)

Professor(Fname, Lname, Dept, Rank, Salary, Age)

Take(Sno, Cno, Grade)

Teach(Fname, Lname, Cno)

Figure 2.1: Example database.

In the tables, the column name abbreviations have the following meaning:

Sno: student number;

Sname: student name;

Cno: course number;

Fname: professor's first name;

Lname: professor's last name;

Dept professor's department.

Relational keys are identified by underlining. Note that table Professor has a composite- two column-key. We have split

professors' names into separate first and last names to motivate certain SQL features introduced later .

11 Example Datablle

11 THE ESSENCE OF SQL: A Gulde to Learning

Most of SQL In the Least Amount of Time

3 A Bird's Eye View of SQL

SQL can be divided into four parts, or sublanguages, as they are traditionally called: data definition language (DDL), data manipulation language (DML), system administration language (SAL) and query language.

DDL provides facilities for creating and destroying tables and indices, as well as for affecting their physical layout. A simple example of DDL, creating our table Student and its related index is shown in Figure 3.1. (In this essay, we use generic SQL syntax which, except for the data type names and minor syntactic differences, is applicable to all SQL dialects.)

```
CREATE TABLE Student(  
    Sno char(9) Sname NULL)  
    varchar(20) Age int  
    NOT NULL, NULL,  
  
CREATE UNIQUE INDEX Student_ndx ON  
    Student(Sno)
```

Rgure 3.1: DDL for table Student.

The table creation syntax closely mimics record layout definitions in most modern 3GLs. Concept of NULL and the

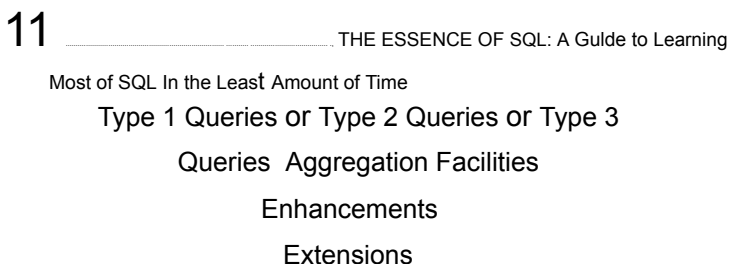


Figure 3.2: SQL query types.

Types 1, 2 and 3 are the three distinct query types, each motivated by a different category of questions and database designs. For those already familiar with SQL, Type 1 queries correspond to "flat" (single level) queries, and Type 2 and Type 3 queries correspond to "nested" (multi-level) queries: Type 2- without correlations, and Type 3-- with correlations.

Aggregation facilities, enhancements and extensions are orthogonal to these three query types, and can be used in conjunction with all of them. They can also be used in combinations with each other.

Aggregation facilities provide means for computing various aggregates, such as sums and averages, and for grouping the data.

Enhancements are those features that make SQL more practical, but do not fundamentally increase its expressive power. answer is one such
For example, ability to sort the
enhancement.

Extensions encompass various advanced features- most important of which is the explicit WHILE loop- that have been added to SQL recently by various vendors. These

extensions make the language fundamentally more powerful; however, they are still quite non-standard, and are considered only briefly at the end of this essay.

11 THE ESSENCE OF SQL: A Guide to
Learning Most of SQL In the Least Amount of Time

4 The Standard Questions

We begin our presentation of SQL with a list of standard

questions, as shown in Figure 4.1. (Designation EI stands for "English question I," etc.) This list will be extended with additional standard questions as we proceed. The key subphrases that make these questions standard are italicized for emphasis.

- EI. Who takes (the course with the course number) CS112?
(By "Who" we mean that we want the student numbers retrieved. If we want the names retrieved, we will explicitly say so.)
- E2. What *are* student numbers and names of students who take CS112?
- E3. Who takes CS112 or CS114?
- E4. Who takes both CS112 and CS114?
- ES. Who does not take CS112?
- E6. Who takes a course which is not CS112?
- E7. Who takes at least 2 courses (i.e., at least 2 courses with different course numbers)?
(A more general question: Who takes at least 3, 4, 5, etc., courses?)
- ES. Who takes at most 2 courses?
(More generally: 3, 4, 5, etc., courses?)
- E9. Who takes exactly 2 courses?
(More generally: 3, 4, 5, etc., courses?)
- EIO. Who takes only CS112?
- EII. Who takes either CS112 or CS114?
- EI2. Who are the youngest students?
(Similarly, Who are the oldest students?)
- EI3. Who takes every course?

Figure 4.1: A list of standard questions.



5 The First Two Questions

The first two questions fall into the category of Type 1 queries. All Type 1 SQL queries have the basic form

```
SELECT <list of desired columns>
FROM <list of tables>
WHERE      condition>
<Boolean
```

where SELECT, FROM and WHERE are keywords designating the three basic components, or clauses, of the query. (While SQL is generally case insensitive, for clarity we will show all keywords in upper case.)

In thinking about, understanding and composing SQL queries, we look at the FROM clause first. This clause lists the table(s) that need to be considered to answer the question. (To determine this list, pretend that you have to answer this question without a computer system, using just paper copies of the tables, and think of which of them you would actually need.)

In question E1- Who takes CS112?- "Who" stands for student numbers (Sno), and CS112 stands for a course number (Cno). Thus, since all relevant information is contained in table Take, the FROM clause becomes

```
FROM Take
```

Next, we consider the WHERE clause. It contains a Boolean condition which defines which rows from table(s) in the FROM clause should be retrieved by the query. For question E1 this condition is for the course number to be CS112, thus making the WHERE clause

```
WHERE (Cno = "CS112")
```



(Whether one uses single or double quotes around string literals is usually system dependent. Also, while parentheses around conditions are often not required, we use them to improve readability.)

Finally, the SELECT clause lists the column(s) that define the structure of the answer. Since in this case we just want the student numbers retrieved, the SELECT clause becomes

```
SELECT Sno
```

The final query then takes the form shown in Figure 5.1.

```
SELECT Sno
FROM      Take
WHERE     "CS112")
(Cno =
```

Figure 5.1: Query Q1 for question E1 Who takes CS112?

While it is true that this query does intuitively correspond to question E1, intuition is not a reliable means for understanding SQL. So, to be safe in interpreting SQL queries, we introduce a conceptual device known as the query evaluation mechanism.

A query evaluation mechanism gives us a precise and formal algorithm to trace queries.

This ability to trace is fundamental to all SQL programming, since it is only by following the trace that we can see how the answer is actually computed, and thus understand its true meaning.

Different types of SQL queries have different evaluation mechanisms. The Type 1 evaluation mechanism is shown in Figure 5.2

-
1. Take a cross-product of all tables in the FROM clause - i.e., create a temporary table consisting of all possible combinations of rows from all of the tables in the FROM clause. (If the FROM clause contains a single table, skip the cross-product and just use the table itself.)
 2. Consider every row from the result of Step 1 exactly once, and evaluate the WHERE clause condition for it.
 3. If the condition returns True in Step 2, formulate a resulting row according to the SELECT clause, and retrieve it.

Figure 5.2: The Type 1 SQL evaluation mechanism.

Given this evaluation mechanism, we can now formally trace query Q1, as follows.

1. Since query Q1 involves only one table, we effectively skip Step 1.
2. We look at every row from table Take, and evaluate condition (Cno = "CS112") for it.
3. We take the Sno values from those rows where this condition returns True, and place them into the result.

By following this trace, we can now confidently assert that query Q1 indeed corresponds to our question E1. (Do not be deceived into thinking that just because this evaluation mechanism and the consequent trace are so simple, using the mechanism is the same as just using the basic intuition. As we will demonstrate very shortly, this is not at all so.)

11 The First Two Questions

We note that because evaluation mechanisms are formal devices, they do not represent actual evaluation strategies taken by real systems. (No real system would be caught dead actually taking cross-products all the time.) All that is required of them is that they always give the same result as a real

system. Thus, in developing this particular form of the evaluation mechanism, we were not concerned with its apparent inefficiency, and concentrated instead on its clarity and ease of use in tracing.

The SQL query corresponding to question E2- What are student numbers and names of students who take CS112?- is shown in Figure 5.3.

```
SELECT Student.Sno, Sname
FROM Student, Take
WHERE
    (Student.Sno = Take.Sno)
AND
    "CS112")
(Cno =
```

Figure 5.3: Query Q2 for question E2. What are student numbers and names of students who take CS112?

This query shows an example of using several tables in the FROM clause and several columns in the SELECT clause. It also introduces a new syntactic feature.

Because cross-products retain all columns from all of the tables involved, in this case the result of the cross-product will have two columns labeled Sno—one from the Student table and the other from the Take table. Thus, every time we refer to Sno, we need to explicitly specify, or disambiguate, which of the two Sno columns we mean. The use of the "dotted" notation Student.Sno and Take.Sno achieves this. This notation (which is quite standard in most programming languages for use with record variables and their





fields) is called a prefix notation in SQL, with "Student." and "Take." called prefixes.

Also, while it does not matter which of the two Sno columns- Student.Sno or Take.Sno- we choose for the answer (after all, they are equal to each other), we must explicitly specify one of them in the SELECT clause. Not doing so would cause a syntax error. (An informal explanation for this is as follows: Column name disambiguation is a syntactic issue and must be resolved at compile-time, while equality is not known until run-time.)

Column name disambiguation is a syntactic issue and must be resolved at compile-time, while equality is not known until run-time.)

We now trace this query, using our evaluation mechanism, as follows. (To better follow this trace, make some example data for the tables, and execute its steps on paper.)

1. We take the cross-product of tables Student and Take. This cross product would contain every combination of rows from these tables. (A good way to visualize the result of this cross-product is to think of it as comprised of "wide" rows formed by "concatenating" the Student and Take rows from each combination.)
2. We evaluate the WHERE clause condition for every such combination. This condition has two parts, with the conjunct (Student.Sno = Take.Sno) assuring that the Student row and the Take row in the combination deal with the same student, and the conjunct (Cno = "CS112") assuring that we are dealing with course CS112.
3. For those combinations where the condition returns True, we choose the Sno and

Sname values from the Student row and put them into the answer.


Again, it is by following this trace, that we can confidently conclude that query Q2 indeed corresponds to question E2.

While required to disambiguate which columns come from which tables when there are several same-named columns in the result of the cross-product, the prefix notation is always permitted for all column references. Thus, even though it is not necessary to use a prefix with Sname or Cno—there is only one column called Sname and only one column called Cno in the cross-product, it is perfectly legal to write our query Q2 as shown in Figure 5.4.

```
SELECT Student.Sno, Student.Sname
FROM
    Student, Take
WHERE      (Take.Cno =
(Student.Sno =  "CS112")
Take.Sno)
AND
```

Figure 6.4: Query Q2 rewritten with all explicit prefixes.

The essential feature of the Type 1 evaluation mechanism is that every wide row from the result of the cross-product is considered exactly once. Thus, even though the order in which these rows are considered is not specified, fundamentally, Type 1 queries involve only one pass through the data.

This means that the decision of whether or not some row combination from the cross-product will contribute to the answer has to be made exactly when this combination is considered, and not at any other time during the evaluation. In particular, it cannot be delayed until after some other row combinations have been looked at. As we will see shortly, this one-pass-only property of Type 1 queries  be of great

6 Standard Questions Involving "or" and "both-and"

Our standard question E3- Who takes CS112 or CS114?- is posed by query Q3 shown in Figure 6.1.

```
SELECT Sno
FROM Take
WHERE      "CS112")
(Cno =
      OR (Cno ="CS114")
```

Figure 6.1: Query Q3 for question E3 - Who takes CS112 or CS114?

A trace of this query shows its correctness. The only new issue here is this: If some student actually does take both CS112 and CS114, then his Sno will be retrieved twice by the query- the first time when the query processes a Take row with his Sno and CS112, and the second time when it processes a Take row with his Sno and CS114. Such duplicates can be eliminated by rewriting the SELECT clause as follows:

```
SELECT DISTINCT Sno
```

The keyword DISTINCT, which eliminates duplicate rows from the answer, is one of the enhancements available in SQL. (It is considered an enhancement because it does not change the meaning of the answer- it simply makes it more concise.) However, its use is expensive, since duplicate elimination requires the answer to be sorted. Thus, it should be used only when concerns for the clarity of the answer outweigh those for its efficiency.

An interesting variation of this query is to rewrite the WHERE clause as

```
WHERE (Cno IN ("CS112", "CS114"))
```

which uses the list membership operator IN. (Operator IN returns True if the value on its left is equal to one of the values in the list on its right.) While just providing an alternative formulation in this case, this operator is actually a significant feature of SQL and will become necessary when we introduce Type 2 queries.

Our standard question E4- Who take both CS112 and CS114?- presents a much more interesting case. First, consider the query of Figure 6.2.

```
SELECT Sno
FROM Take
WHERE      "CS112")
(Cno =
  AND (Cno = "CS114")
```

Figure 8.2: First attempt at question E4 Who takes both CS112 and CS114?

This query, which was generated from query Q3 by replacing OR with AND, reflects a misplaced intuition that logical operators directly model, and thus can be automatically substituted for, their English counterparts. However, intuition is not a reliable tool when programming in SQL, and the query of Figure 6.2 is not a correct implementation of question E4. It

will compile, however, and therefore will generate some answer. (Before proceeding further, try to figure out exactly what answer this query will generate, and why it is incorrect.)

11 THE ESSENCE OF SQL: A Guide to

Learning Most of SQL In the Least Amount of time

Recall that the conditions of Type 1 queries are evaluated on a row-by-row basis, once per row. Since no row can have a Cno value which is simultaneously equal to "CS112" and to "CS114", the condition of this query will always return False. Thus, the answer to this query will always be empty.

While question E4 can be posed as a Type 1 query, to discover this solution, we need to employ a bit of "reverse engineering." In other words, we first need to figure out how we can actually compute the answer to this question in a manner consistent with the Type 1 evaluation mechanism- in effect, visualizing the corresponding trace, and then to "back into" the SQL query itself. We note that:

This ability to visualize traces and then reverse engineer the appropriate SQL code is the key to one's mastery of SQL.

To visualize the trace, we first need to determine which table(s) need to be considered. Since the question Who takes both CS112 and CS114? involves only student numbers (Sno) and course numbers (Cno), the only table that is needed is table Take.

However, as we have just argued, since every Take row contains only one Cno value, and since it is considered only once by the evaluation mechanism, we cannot evaluate our condition directly on the rows of Take. What we need instead is a table with rows that for each student would list not one, but two, course numbers for the courses he takes.

But how can we generate such a table from our basic table Take? The answer is: In two steps. First, we can compute a cross-product of table Take with itself. (Think of making a copy

11 Standard Questions Involving "or" and

"both.and"

Second, we can impose equality on the two Sno values in each wide row in the result of this cross-product. This would remove those wide rows where the Sno value from the first copy of Take is different from the Sno value from the second copy.

Each remaining wide row would then refer to just a single student, and the two- one from each copy of Take- Cno values for the courses he takes. We can then test one of them to be "CS112" and the other to be "CS114," and for those rows where these tests would succeed, retrieve the Sno value into the answer.

Since in Type 1 SQL queries cross-products are caused by listing tables in the FROM clause, we can informally sketch a query that would have such a behavior as shown in Figure 6.3.

```
SELECT Take.Sno
FROM Take, Take
WHERE      Take.Sno)
(Take.Sno =
    AND (Take.Cno = "CS112")
    AND (Take.Cno = "CS114")
```

Figure 6.3: A sketch of the query for E4 .Who takes both CS112 and CS114?

There is an obvious problem with this sketch, however. Since both tables in the FROM clause have the same name (one of them, after all, is a copy of the other), the use of table names for prefixes is not sufficient to disambiguate column references.

What we really want to say here is that in the condition (Take.Sno = Take.Sno) the left Take.Sno is taken from the

first copy of Take, and the right Take.Sno is taken from the second copy of Take. Similarly, in the condition (Take.Cno = "CS112") we mean the first copy of Take and in the condition (Take.Cno = "CS114") we mean the second one. Finally, since both Sno values are the same in the wide row, it does not matter to us which one is retrieved in the SELECT clause, but we must explicitly choose one or the other. None of this, however, is reflected in our query sketch.

To take care of such cases, SQL allows any table in the FROM clause to be optionally followed by its temporary **alias**-a new table name, which is valid only for the duration of the particular query and which is used for prefixes. The aliasing feature allows us to phrase the correct query for question E4 as shown in Figure 6.4.

```
SELECT X.Sno
FROM Take X, Take
WHERE AND (X.Cno = "CS112")
AND      (Take.Cno =
(X.Sno = Take.Sno)"CS114")
```

Figure 6.4: Query Q4 for question E4 .Who takes both CS112 and CS114?

Before proceeding further, we want to make several comments regarding the use of aliases. First, the choice of alias names is completely arbitrary as long as they do not conflict with each other or with any real table name used in a query.

Second, aliases are **opaque**- a technical term meaning that an alias completely covers and hides the name of the underlying table. Thus, from the point of view of the SELECT and WHERE clauses, the query of Figure 6.4 involves two tables: one named X and the other named Take.

11

Standard Questions Involving "or" and
"both/and"

Third, it is always permitted to alias tables- even if not really

necessary. Indeed, to save on typing, SQL programmers often use short aliases for long, meaningful table names given to them by well-meaning database designers. Consequently, our query Q4 could also have been written as shown in Figure 6.5. We stress, however, that using prefix "Take." anywhere in this query would now be syntactically incorrect.

```
SELECT X.Sno
FROM Take X, Take Y
WHERE      Y.Sno)
AND        (X.Cno =
(X.Sno =   "CS112")
      AND (Y.Cno = "CS114")
```

Figure 6.5: Q4 rewritten with both copies of Take aliased.

Finally, we note that aliases are an essential feature of SQL. Without them, standard questions involving the both-and construct could not have been asked as Type 1 queries.

7 Negation in SQL

Negation is one of the most interesting and complex issues in SQL. Both question E5-Who does not take CS112?, and question E6-Who takes a course which is not CS112? involve negation.

It is important to understand that these are not restatements of each other, but are indeed different questions. The following two example situations, for which these two questions would generate different answers, illustrate this point. (Try to find them before proceeding.)

The first example involves a student who takes both CS112 and some other course, say CS113. Because this student takes CS112, he will not be in the answer to question E5. However, because he takes a course-CS113, in this case-which is not CS112, he will be in the answer to question E6.

The second example involves a student who does not take any courses. (Or, to put it another way, who takes nothing.) Since he does not take any courses, it follows that he does not take CS112; therefore, he will be in the answer to question E5. However, he will not be in the answer to question E6, which requires him to take some course. (Note that the student number of this student would be included in table Student, but would be absent from table Take. For those familiar with the concepts of referential integrity and foreign key constraints, this situation is perfectly legal and would be allowed in this database.)

11 Negation in SQL

Not only are questions E5 and E6 different, they are orthogonal to each other- i.e., all four combinations of

Yes/No answers for these two questions are possible, as shown in Figure 7.1.



Situation Included in answer to E6
Included in answer to ES

Student takes both CS112 and CS113 No Yes

Student takes nothing Yes No

Student takes only CS112 No No

Student takes only CS113 Yes Yes

Figure 7.1: Orthogonality of questions ES and E6.

The query posing question E6- Who takes a course which is not CS112?-is shown in Figure 7.2.

```
SELECT Sno
FROM Take
WHERE      "CS112")
(Cno !=
```

Figure 7.2: Query Q6 for question E6-Who takes a course which is not CS112?

A trace of this query shows its correctness. We only note that the answer to this query may contain duplicates. To eliminate them we would need DISTINCT in the SELECT clause.

This is a good time to mention that in addition to equality (=)

and inequality (which in some SQL dialects is expressed as \neq and in some as \neq), SQL supports all other standard binary comparators: $<$, \leq , $>$, and \geq . All SQL dialects support these latter comparators for numeric data types; most also support them for strings and other data types as well.

Question ES- Who does not take CS112?-is much trickier. First, note that rewriting the WHERE clause as

WHERE NOT (Cno = "CS112")

will not work, because conditions NOT (Cno = "CS112") and (Cno \neq "CS112") are equivalent to each other, and thus would both pose question E6. (We note that SQL supports full Boolean logic including DeMorgan's and standard Boolean Distributive and Associative Laws.)

A more fundamental problem with question ES emerges when one considers the basic one-pass nature of Type 1 queries. As it turns out, the answer to this question cannot be computed in a single pass, even if we do take a cross product of tables Student and Take. (Try it on some data examples, and observe that one pass is not enough.)

What is required, instead, are two passes: the first pass through table Take to select those students who do take CS112, and the second pass through table Student to "get rid of" them. Furthermore, the first pass has to be completed - i.e., we have to identify and in a way "collect" all of these "bad students," before starting on the second pass to eliminate them.

11 Negation In SQL

Since the Type 1 SQL evaluation mechanism is fundamentally "one pass," it is impossible to pose question E5 as a Type 1 query, given this or any other reasonable database design.

Instead, question E5 requires subqueries, which are provided by Type 2 and 3 queries, and which are covered later.

An important general difference between questions E5 and E6 is that E5 involves "real negation" and E6 involves the so called "pseudo-negation." Questions involving pseudo negation can be posed as Type 1 queries; questions involving real negation cannot.

One way to distinguish between the two is to see what is being negated- some noun constant (e.g., "is not course CS112" in E6) or a verb (e.g., "does not take" in E5). Also, question Q6 can be rephrased by replacing the English word "not" with the phrase "different from" (or "other than"), as in Who takes some course different from CS112? This technique is a rule of thumb that works most of the time.

least

Question E7- Who takes at least 2 courses?- is the last of our standard questions that can be posed as a Type 1 query. The query corresponding to this question, shown in Figure 8.1, is quite similar to the "both-and" query Q4, and has a very similar trace. (An alternative way to pose this query is to replace the inequality in the WHERE clause by the less than comparison (X.Cno < Take.Cno), assuming that the particular SQL dialect allows less-than comparisons between strings.) To eliminate duplicates from the answer we would need DISTINCT in the SELECT clause.

```
SELECT X.Sno
FROM Take X, Take
WHERE AND (X.Cno !=
(X.Sno =      Take.Cno)
Take.Sno)
```

Figure 8.1: Query Q7 for question E7 Who takes at least 2 courses?

Of course, as noted in our standard questions list, there are many cases of the use of "at least." So, how would we pose the question Who takes at least 3 courses? The answer is: The same way, only using 3 copies of table Take, as shown in Figure 8.2.

```
SELECT X.Sno
FROM Take X, Take Y, Take
WHERE AND (X.Cno !=
AND      Y.Cno) (Y.Cno
AND      != Take.Cno)
AND      (X.Cno !=
(X.Sno =      Take.Cno)
Y.Sno) (Y.Sno
= Take.Sno)
```

Figure 8.2: The query for question: Who takes at least 3 courses?

Again a trace shows the correctness of this query. We note that, because equality is transitive, two equalities on Sno are sufficient. However, because inequality is not transitive, three inequalities on Cno are necessary. We also note that the use of DISTINCT is again necessary to eliminate duplicates.

An interesting variation here is to replace the three inequalities among the Cno's in the WHERE clause by the condition $(X.Cno < Y.Cno) \text{ AND } (Y.Cno < Take.Cno)$. First, since less-than comparisons are transitive, only two of them are necessary. Second, this also substantially reduces the number of potential duplicates in the answer.

We now have a general approach that gives us a family of Type 1 solutions for "at least K" questions for any K. For "at least 4 courses" we would use 4 copies of Take, 3 equalities on Sno, and either 6 inequalities or 3 less-than comparisons among Cno's. For "at least 5 courses" we would use 5 copies of Take, 4 equalities and either 10 inequalities or 4 less-thans, etc. (The formula for computing the number of required inequalities is $K*(K-1)/2$.)

Clearly, these solutions become very bulky and inefficient for all but the smallest values of K. We will show later how the questions involving "at least" can also be posed in a more concise and efficient way using aggregation. Nonetheless, the ability to pose them as Type 1 queries is important.

9 Negation Revisited

We now consider the remaining questions from our standard list: questions EB through EI3, as well as question ES. As we have already noted, question ES involves real negation—the kind that cannot be phrased as a Type 1 query. Although not obvious at the first glance, these remaining questions also involve real negation—to see that, consider rephrasing them as shown in Figure 9.1. (Negation is italicized for emphasis.)

While the rephrasings for questions E12 and E13 are somewhat obscure (they are, nevertheless, correct), the rest of the rephrasings are quite intuitive. Thus, we have essentially reached a dead-end—to proceed further we need real negation, and Type 1 SQL cannot give it to us. Not to worry, however—that's why we have the rest of SQL.

EB. Who takes at most 2 courses?

is equivalent to

Who does not take at least 3 courses?

E9. Who takes exactly 2 courses?

is equivalent to

Who takes at least 2 courses and does not
take at least 3 courses?

EIO. Who takes only CS112?

is equivalent to

Who takes CS112 and does not take
any other course?

EII. Who takes either CS112 or CS114?

is equivalent to

Who takes CS112 or CS114, and does not
take both CS112 and CS114?

10 Type 2 SQL Queries

A sample Type 2 query is shown in Figure 10.1. We begin the discussion by first explaining its syntax, then introducing the Type 2 evaluation mechanism, and finally tracing this query to determine the question it poses.

```
SELECT Sno, Sname
FROM Student
WHERE Sno FROM
(Sno IN Take
(SELECT
    WHERE "CS112"))))
(Cno =
```

Figure 10.1: A sample Type 2 SQL query.

Syntactically, a Type 2 SQL query is a collection of several non-correlated component queries, with some of them nested in the WHERE clauses of the others. (The meaning of the term "non-correlated" will be explained in a moment.)

Theoretically, there is no limit on the number of nesting levels or on the number of component queries involved. On a practical level, however, SQL compilers do impose some limitations in this regard (e.g., no more than 16 nesting levels or 256 component queries), but these rarely present any real problems.

In the above example, we have two component queries. The inner query (called the subquery) is nested in the WHERE clause of the outer query (called the main query), and the queries are connected by the list-membership operator IN. (Out of the three pairs of parenthesis used in this example, _____

11

Type 2 SQL Queries

the only pair that is actually required is the left parenthesis immediately before SELECT and its matching right parenthesis.)

A query is called non-correlated if all column references in it are local- i.e., all columns come from, or are bound to, the tables in the local FROM clause. Any non-local column reference is called a correlation, and the query containing it becomes correlated.

In determining these column-to-table bindings, SQL follows the standard "inside-out, try the local scope first" scope rules. Specifically, given a column reference, SQL first tries to find some table in the local FROM clause containing that column. If the column reference also involves a prefix either a table name or an alias- then SQL also looks for the match on that prefix. Three alternative outcomes are then possible.

1. SQL finds a single such table and successfully binds that column reference to that table.
2. SQL finds several such tables. Then column reference is ambiguous and the appropriate error message is generated.
3. SQL does not find any such table in the local FROM clause. Then column reference is not local. SQL then looks at the next outer scope-i.e., at the FROM clause in the immediately enclosing query, and attempts to bind that column reference to a table in that FROM clause.

This process continues (using the same three possibilities)

until either a column reference is successfully bound, or an ambiguity is found, or the binding process "falls off" the main query, in which case the column reference cannot be bound at all, and the appropriate error is declared. (This

last case corresponds to an undeclared variable in conventional programming languages.)

Given this process and the query of Figure 10.1, columns Sno and Cno in its subquery are bound to the inner Take table, and columns Sno and Sname in its main query are bound to the outer Student table, thus making all bindings local, the component queries non-correlated, and this entire query of Type 2.

The formal syntactic condition for a multi-level (i.e., with subqueries) query to be of Type 2 is presented in Figure 10.2.

A query with subqueries is of Type 2 if every component query in it is non-correlated or, equivalently, if every column reference in it is local.

Figure 10.2: The definition of a Type 2 query.

The evaluation mechanism for Type 2 queries is presented in Figure 10.3.

To execute a Type 2 query, execute its component queries in the "inside-out" order- i.e., with the inner-most nested subquery first, replacing each query by its result as it gets evaluated.

Figure 10.3: The Type 2 SQL evaluation mechanism.

In case of the query of Figure 10.1, the subquery

```
SELECT Sno  
FROM Take  
WHERE (Cno = "CS112")
```

. .



which is a regular Type 1 query, is executed first. Its answer is then substituted into its place in the main query

```
SELECT Sno, Sname  
FROM Student  
WHERE (Sno IN ( ... ))
```

which is executed next. Note that, at this point, the main query has been reduced to a simple Type 1 query. Also note that, since the subquery retrieves a single column in its SELECT clause, the answer to it is just a list of values. Thus, the use of the list membership operator IN as the inter query connector is quite appropriate.

We note that Type 2 queries fundamentally involve multiple data passes - one for each component query. In this case, the first pass is through table Take in the subquery, and the second pass is through table Student in the main query.

To determine the question posed by this query, observe that the subquery is a verbatim copy of query Q1, and thus poses question E1- Who takes CS112? The main query, which retrieves their student numbers and names then cor

responds to the question What are the student numbers and names of students who take CS112? In other words, this is just another way of asking our standard question E2.

Before concluding this section, we note that, from a

syntactic point of view, the condition (Sno IN ...) of the main query is just that- a condition; thus, it can itself be part of a more complex Boolean expression involving NOT, AND and OR- a feature that will become very handy in a moment.

111 _____ THE ESSENCE OF SQL: A Guide to Learning Most of SQL In the Least Amount of Time

11 Using Type 2 to Implement Real Negation

As we have discussed in Section 7, questions involving real negation need two passes through the data, using the following general strategy:

- To pose a question "Who does not do X? "
1. identify and select those who actually do X; and
 2. remove them from the list of those who potentially may do X.

Since Type 2 queries fundamentally involve multiple data passes, they give us exactly what is necessary to implement real negation in SQL. Consider the query of Figure 11.1.

```
SELECT Sno
FROM Student
WHERE      FROM Take
NOT (Sno IN

(SELECT Sno
      WHERE      "CS112"))
(Cno =
```

Figure 11.1: A mystery Type 2 SQL query Involving negation.

This query differs from the query of Figure 10.1 in two ways: there is a NOT in front of the main query's condition, and Sno is used alone in the main SELECT clause. But how is this query evaluated, and what question does it pose?

Since this is a Type 2 query, it is evaluated inside-out. Again, the subquery is the same as query Q_t and poses question E_I- Who takes CS112? Because of the NOT operator in its WHERE clause, the main query now retrieves the

11 Using Type 2 to Implement Real Negation

student numbers of all other students- i.e., those student numbers that are not on the list generated by the subquery. Thus, the full query of Figure 11.1 corresponds to the question Who does not take CS112?- i.e., our standard question E_S. (We will refer to this query as Q_S.)

An interesting related query is shown in Figure 11.2.

```
SELECT Sno
FROM Take
WHERE      FROM Take
NOT (Sno IN
      (SELECT Sno
        WHERE      "CS112"))
      (Cno =
```

Figure 11.2: Using Take In both FROM clauses.

This query was obtained from the query of Figure 11.1 by replacing table Student with table Take in the main FROM clause.

Note that this is still a Type 2 query, with the inner Sno and Cno coming from the inner Take table, and the outer Sno coming from the outer Take table. (Even though we use two

copies of table Take here, because of the scope rules, no binding ambiguities arise and no aliases are necessary.)

The use of table Take in the main FROM clause limits the list of students who may potentially appear in the answer to those who are listed in Take and thus take some course. Thus, it corresponds to the question Who takes some (i.e., at least 1) course, but does not take CS112?

As we can see, this query is different from the query of Figure 11.1, just as this question is different from our original question ES-Who does not take CS112?

11 _____ THE ESSENCE OF SQL: A Guide to Learning Most of SQL in the Least Amount of Time

12 Posing Questions Involving "at most," "exactly," "only," and "either-or"

Using the solution to question ES as a foundation, it is now quite easy to generate solutions to questions ES through E11. Specifically, question ES- Who takes at most 2 courses?- is posed by query QB of Figure 12.1.

```
SELECT Sno
FROM Student
WHERE
    NOT (Sno IN
        (SELECT X.Sno
         FROM
             Take X, Take Y, Take
          WHERE (Y.Sno =
                (X.Sno = Take.Sno)
                Y.Sno) (X.Cno !=
                AND AND Y.Cno)
                AND (Y.Cno !=
```

Take.Cno)
AND (X.Cno != Take.Cno)))

Figure 12.1: Query QB for question ES .Who takes at most 2 courses?

This query is based on rephrasing question E8 as Who does not take at least 3 courses? The subquery here poses the question Who takes at least 3 courses? and is taken from Figure 8.2. The NOT in the main WHERE clause achieves the desired negation.

We note that this query will retrieve students who do not take any courses. This is appropriate, since "at most 2" means 0 (!), 1 or 2.

By using Take instead of Student in the main FROM clause we can change this query to ask a related question-Who takes some (i.e., at least 1), but at most 2, courses? In other words, Who takes 1 or 2 courses?

Question E9- Who takes exactly 2 courses?- is posed by query Q9 of Figure 12.2.

```
SELECT X.Sno
FROM Take X, Take
WHERE
    (X.Sno = Take.Sno)
    AND      Take.Cno)
    (X.Cno <
    AND NOT (X.Sno IN
              (SELECT X.Sno
               FROM
                 Take X, Take Y, Take
                WHERE   AND
```

```


(X.Sno = Y.Sno)
(Y.Sno = Take.Sno)
AND
(X.Cno < Y.Cno)
AND
(Y.Cno < Take.Cno)))

```

Figure 12.2: Query Q9 for question E9 Who takes exactly 2 courses?

This query is based on rephrasing question E9 as Who takes at least 2 courses and does not take at least 3 courses? The main query poses Who takes at least 2 courses? and is copied from query Q7. (We have used the less-than operator in its condition to remove duplicates from the final answer.) The subquery again poses Who takes at least 3 courses? (Here, we have used less-than operators for conciseness.) The NOT in the main WHERE clause again achieves the desired negation.

We note that even though we have used the same alias name X in both the main query and in the subquery, because of the scope rules, no confusion arises.



Question E10- Who takes only CS112?- is posed by query Q10 of Figure 12.3.

```

SELECT Sno
FROM Take
WHERE NOT (Sno IN
(Cno =
"CS112")
AND
(SELECT Sno
FROM
Take
WHERE (Cno !=

```

"CS112"))))

Figure 12.3: Query Q10 for question E10 -Who takes only CS112?

This query is based on rephrasing question E9 as Who takes CS112 and does not take any other course? Here, the main query is based on query Q1- Who takes CS112? The sub query is taken verbatim from query Q6-Who takes a course which is not CS112? The NOT in the main WHERE clause again achieves the desired negation.

Two things should be noted about query Q10. First, thinking that the outer NOT can be brought inside the WHERE clause of the subquery, and thus cancel the negation implied by the != operator, is a common mistake- it can not. Doing so would change the meaning of the query.

Second, the condition (Cno = "CS112") in the main WHERE clause is unnecessary in this case (!) and can be dropped without changing the query's meaning. The simplified query is shown in Figure 12.4. (Before proceeding further, try tracing it to determine why this is so.)

....., : ? ? ? : : , ? ? : , ? ? ?
 ? : ? ? ? : : : : ? ? : : : 11

```
SELECT
Sno FROM
Take
WHERE      FROM Take
NOT (Sno IN

(SELECT Sno
      WHERE      "CS112"))
(Cno !=
```

For a particular Sno to appear in the final result, it must not appear with any Cno value different from "CS112" in table Take. (If it did, the NOT ... IN of the main condition would "remove" that Sno from the answer.) But, for this Sno to even be considered in the main query, it must come from some Take row, and, thus must have some Cno associated with it in that row. Since this Cno cannot be anything other than CS112, it must be CS112. Thus, explicitly testing for it is unnecessary.

Finally, question E11- Who takes either CS112 or CS114?- is posed by query Q11 of Figure 12.5.

```
SELECT Sno
FROM Take
WHERE                               (Cno = 'CS114')
((Cno = "CS112") OR
  AND NOT (Sno IN
            (SELECT X.Sno
             FROM
               Take X, Take
             WHERE      Take.Sno)
            (X.Sno =
              AND (X.Cno = "CS112")
              AND
                (Take.Cno = "CS114"))))
```

Figure 12.5: Query Q11 for question E11 - Who takes either CS112 or CS114?

11- _____ 'THE ESSENCE OF SQL: A Guide to Learning Most of SQL In the Least Amount of Time

This query which is based on rephrasing question E11 as Who takes CS112 or CS114, and does not take both CS112 and CS114? is a combination of query Q3- Who takes

CS112 or CS114?- for the main query, and query Q4-Who takes both CS112 and CS114?- for the subquery, with the outer NOT achieving the negation.

Note the extra pair of parenthesis around the OR in the main condition. This forces the OR to be executed before the AND in the main condition- otherwise, the SQL standard rules of precedence for logical operators are: first NOT, then AND, and only lastly OR.

13 Computing Extremes

as Type 2 Queries with Negation

Consider question E12 Who are the youngest students ? While the most natural way of posing this question involves the use of the aggregate function MIN() (to be presented later), this question can also be posed using a Type 2 query with negation, as shown in Figure 13.1.

```
SELECT Sno
FROM Student
WHERE
NOT (Age IN
      (SELECT X.Age      Student X, Student
       FROM
            WHERE      Student.Age)))
      (X.Age >
```

Figure 13.1: Query Q12 for question E12 Who are the youngest students?

This solution is interesting because it looks at the problem of minimums (and maximums) in a novel way. It is also important, because, as we will point out later, SQL imposes certain limitations of the use of its aggregate functions. Thus, ability to ask this and similar questions without the use of aggregates is important.

To see how this query operates, let's trace its execution. (You should follow this trace on paper.) Since this is a Type 2 query, it is evaluated inside-out. That means that the subquery is evaluated first. But, what does this subquery retrieve?

Negation

It retrieves a particular Age value if there is some other Age value less than it. In other words, it retrieves those Age values which are not smallest. To rephrase this even shorter, it retrieves all Ages except the smallest one.

Given each student, the main query then tests whether his age is not in the list of ages retrieved by the subquery. Of course, since the only age not present in the list is the smallest one, only the students possessing it- i.e., the youngest ones- will be retrieved in the final answer.

If all students happen to be of the same age, then all students would be youngest. In that case, the answer to the subquery would be empty; the main WHERE clause would be True for all Student rows; and all students would be retrieved.

Note that, in Figure 13.1, the connection between the main query and the subquery is made on the Age column. However, it is also possible to ask the same question by making this connection on the Sno column, by replacing the WHERE clause as follows. (The choice between these two alternatives is a matter of personal preference.)

```
WHERE NOT (Sno IN
            (SELECT X.Sno
             FROM Student X, Student
              WHERE      Student.Age)))
            (X.Age >
```

If we wish to pose the question Who are the oldest students?, we can simply replace the greater-than operator in the clause of the subquery with

the less-than operator.
WHERE

11 _____ THE ESSENCE OF SQL: A Guide to Learning Most of SQL In the Least Amount of Time

14 A Look at SQL Enhancements

In this section, we take a break from the fundamental features of the language and take a brief look at the various enhancements it provides. Roughly speaking, these enhancements can be divided into three categories: syntactic enhancements, presentation enhancements and scalar expressions.

Syntactic enhancements are those features that make the writing of SQL queries easier. They include: omitting the WHERE clause, using symbol * (asterisk) in the SELECT clause, and using operator NOT IN.

SQL permits the omission of the WHERE clause if the condition of the query is always True-- i.e., there really is no condition that needs to be imposed. Thus, a request to retrieve the student numbers of all students is posed by the query of Figure 14.1.

```
SELECT Sno  
FROM Student
```

Figure 14.1: Example of query with the omitted WHERE clause.

SQL provides the shorthand symbol * (asterisk) if the request is to retrieve all columns from some table in the FROM clause. For example, the query of Figure 14.2

retrieves all information about students who take CS112.

..... 11

```
SELECT Student.*
FROM Student, Take
WHERE      Take.Sno)
(Student.Sno =
  AND      "CS112")
(Cno =
```

Figure 14.2: Using symbol * to mean "all columns."

We note that using symbol "*" without the prefix,

as in `SELECT*`

causes all columns from all tables in the FROM clause to be retrieved.

The NOT IN operator is a syntactic enhancement that allows NOT to be brought inside parenthesis containing IN. In other words, any construct of the form

`NOT (<Something> IN (SELECT .. FROM ... WHERE ...))`

can be equivalently rewritten (i.e., without changing its meaning) as

`(<Something> NOT IN (SELECT . . FROM .WHERE ...))`

While NOT IN is actually a primitive operator in SQL, we can think of it simply as a shorthand for the longer NOT ... IN combination, making such combinations easier to

read.

Presentation enhancements make answers easier to understand. They include: suppressing duplicates in the answer, assigning new names to columns of the answer, and sorting the answer.

11 THE ESSENCE OF SQL: A Guide to

Learning Most of SQL In the Least Amount of Time

As previously discussed, using the keyword `DISTINCT`, we can suppress duplicates from the answer.

We emphasize that `DISTINCT` eliminates duplicate rows, and not individual values, from the answer. A simple query that makes this point clear is shown in Figure 14.3.

```
SELECT DISTINCT Sname, Age  
FROM Student
```

Figure 14.3: Suppressing duplicates from the answer.

This query would eliminate an answer row only if it duplicates both the `Sname` value and the `Age` value from some other answer row. Thus, it would leave both rows `<"John Doe", 20>` and `<"John Doe", 23>` in the answer. (Recall that `Sno`, and not `Sname`, is the key for table `Student`; so, there can be two students, each named John Doe.)

Often, names of columns in the `SELECT` clause are not the ones desired for the column headings in the answer. To overcome this, SQL allows for the assignment of new column headings. Depending on the SQL dialect, there are basically two ways to do this. We can prefix the column reference in the `SELECT` clause with its new name followed by the equal sign.

SELECT Name = Sname, ...

Alternatively, we can follow the column reference by the keyword AS and the new name.

SELECT Sname AS Name, ...

11

A Look at SQL Enhancements

In both cases, when we want the new column heading to include spaces (and, in some SQL dialects, certain other characters as well), the new name should be enclosed in quotes.

SELECT "Student Name" = Sname, ...

The sorting of a query answer can involve one or several columns from the answer, and may also specify the direction (ascending or descending) for each of these columns. As an example, consider the query of Figure 14.4.

```
SELECT Sname, Age
FROM Student, Take
WHERE      Take.Sno)
(Student.Sno =
      AND (Cno = "CS112")
ORDER BY Sname, Age DESC
```

Figure 14.4: Sorting facilities of SQL

The sort specification is contained in the ORDER BY clause, which syntactically appears as the very last clause of the query, and which is executed after the SELECT clause.

The argument to the ORDER BY clause is a list of columns from the SELECT clause, where the left-to-right column order defines the major-to-minor sort sequence. Each column can optionally be followed by the keywords ASC or

DESC, defining the sort order on that column as ascending or descending, respectively, with ASC as the default if omitted.

In this case, the answer would first be sorted by Sname ascending, and then, within each group of rows with the same Sname, by Age descending.

11 _____ THE ESSENCE OF SQL: A Guide to

Learning Most of SQL In the Least Amount of Time

An interesting feature of the ORDER BY clause is that it also allows us to refer to columns from the SELECT clause by their position. Furthermore, mixing of column names and positions is allowed. Thus, for example, the ORDER BY clause from the query of Figure 14.4 can be equivalently rewritten as

ORDER BY Sname, 2 DESC

Notably, this is the only place in the entire SQL language where columns can be referenced by their positions. (This feature will become useful in a moment.)

We note that changing the ORDER BY clause in our query to

ORDER BY Age DESC, Sname

only changes how the answer is displayed row-wise- it will now be sorted first by Age descending, and then by Sname. Columns will still be displayed as specified in the SELECT clause-with Sname column on the left, and Age column on the right.

Finally, the behavior of sorts is always consistent with the behavior of the less-than and greater-than operators. For example, if the less-than operator treats strings as left-justified, the ascending sort would place them in the standard lexicographic (dictionary) ordering.

The most important enhancement feature of SQL is the

availability of scalar expressions. Luckily, their use is very intuitive. Basically, SQL provides a reasonable set of scalar operators and functions for operating on numerics, strings and other data types, and allows their use in all reasonable places in the SELECT and WHERE clauses of queries.

11

A Look at SQL Enhancements

While specifics as to which operators and functions are actually provided differ widely among SQL dialects, as an example here is a sampling of what Transact SQL provides.

For numerics, there are the four basic arithmetic operators (+, -, *, and /), modulo division (%), functions round(), trunc() and abs(), exponentiation function power(), logarithmic functions log() and log10() for natural and base-10 logarithms, a full set of trigonometric functions, etc.

For strings, there is a concatenation operator (+) and a set of substring functions, etc.

Other data types (most notably, datetime—for dealing with calendar dates and time values) have rich sets of functions as well.

The basic rule of thumb for using scalar expressions in SQL is as follows: Wherever one can use a column name in the SELECT or WHERE clauses, one can also use an expression of an appropriate data type comprised of column names, constants and various scalar operators and functions.

For example, the query of Figure 14.5 assumes table Room(Rno, Length, Width) and retrieves room numbers and areas of those rooms where the length of the room is within 1% of its width, sorting the answer in the descending order by area.

```
SELECT Rno, Area = Length*Width
FROM Room
```



```
WHERE          dth <= 0.01  
abs(Length-Width)/Wi  
ORDER BY 2 DESC
```

Figure 14.5: Example use of arithmetic in SQL



Since many SQL dialects forbid the use of expressions in the ORDER BY clause, and since in most of them column renaming is done after the sorting, referring to expression based columns by their positions in the ORDER BY clause is the only means left to achieve the desired sort.

A less obvious consequence of having expressions in SQL, is that we can use the string manipulation capabilities for formatting query answers. As an example, Figure 14.6 shows a query that will retrieve all professor names, formatted as last name, comma, space, first initial, period, as in: "Smith, J." (Function `substring(Fname,1,1)` gets one character starting at position 1 from Fname; when used with strings, symbol `+` denotes the concatenation operator.)

```
SELECT Lname + ", " + substring(Fname,1,1) + "." FROM  
Professor
```

Figure 14.6: Formatting an answer.

15 Handling Composite Keys

All of the questions considered thus far in this essay have been about students taking courses. Surely, the same types of questions can also be asked about professors teaching them. Are there any differences, and do the same types of solutions apply?

The answer is: Yes, the same types of solutions apply, but with an adjustment. Specifically, since table Professor has a composite, two column- Fname and Lname- key, any professor-based connection (positive or negative) between tables Professor and Teach, or between multiple copies of Professor or of Teach, has to involve both of these columns.

To illustrate this point, we briefly go through some of our standard questions, restated with respect to professors teaching courses. We begin with questions not involving

negation.

The query of Figure 15.1 poses the question What are full names and ages of professors who teach CS112?, which is similar to our standard question E2.

```
SELECT P.Fname, P.Lname, Age
FROM Professor P, Teach T
WHERE      T.Fname)
(P.Fname =
  AND (P.Lname = T.Lname)
  AND (Cno = "CS112")
```

Figure 15.1: Query for the question What are full names and ages of professors who teach CS112?

11 Handling Composite Keys

Questions about professors teaching courses involving "or," "both-and," and "is not" (similar to standard questions E3, E4, E6) are posed in a similar fashion, and are left as an exercise for the reader.

The question Who teaches at least 2 courses? (similar to question E7) is also quite straightforward, and is posed in a manner similar to query Q7.

Another question also involving "at least"- Which courses are taught by at least 2 professors?- but formulated with respect to courses taught by professors, however, takes more thought. (Try posing it before proceeding.)

The tricky part here is to realize that a difference in just the first names or just the last names of two professors is sufficient to make them different from each other. The correct way of expressing this condition involves the use of OR, as shown in Figure 15.2, with the necessary extra pair of parenthesis around the OR operands. As an additional exercise, try posing the question Which courses are taught by at least 3

professors?

```
SELECT X.Cno
FROM
    Teach X, Teach Y
WHERE    (X.Cno =
AND      Y.Cno)
        ((X.Fname != Y.Fname)
OR (X.Lname != Y.Lname))
```

Figure 15.2: Query for the question Which courses are taught by at least 2 professors?

11 _____ THE ESSENCE OF SQL: A Guide to Learning Most of SQL In the Least Amount of Time

16 Negation Over Composite Keys

Posing negative questions regarding professors teaching courses presents its own set of interesting problems. Consider the question Who does not teach CS112? similar to our question ES- along with the query of Figure 16.1.

```
SELECT Fname, Lname
FROM Professor
WHERE    Fname
(Fname NOTFROM
IN (SELECTTeach
    WHERE    "CS112"))
    (Cno =
AND      IN (SELECT
(Lname NOT Lname
FROM Teach
WHERE    (Cno =
```

Figure 16.1: First attempt at the question Who does not teach CSU2?

However intuitive, this query is nonetheless incorrect for this question. (Try to determine why before proceeding.)

The problem here is that the lists of the first names and the last names of professors who do teach CS112 are compiled and tested independently of each other. Thus, they are not "synchronized" to come from the same Teach row. As a result, professors who do not in fact teach CS112 may incorrectly be suppressed by this query.

To see this consider an example where table Professor contains only two rows, and table Teach contains only a single

11

Negation Over Composite Keys

row, as shown in Figure 16.2. (Since the Dept, Rank, Salary and Age values are immaterial for this example, we are not showing these columns in table Professor.)



Professor

Fname Lname

"John" "Smith"

"Mary" "Smith"

Teach

Fname | Lname | Cno

"John" | "Smith" | "CS112"

Figure 16.2: Example data for tables

Professor and Teach.

A trace of the query of Figure 16.1 on this data shows that it will return an empty result. Yet, Mary Smith does not teach

CS112, and thus should be in the answer to the question Who does not teach CS112?

Changing this query to use OR in place of AND does not fix the problem either. As an example here, consider tables Professor and Teach as shown in Figure 16.3.



Professor	
Fname	Lname
"John"	"Smith"
"Mary"	"Smith"
"John"	"Brown"
"Mary"	"Brown"

11

THE ESSENCE OF SQL: A Guide to
Learning Most of SQL In the Least Amount of Time



Teach		
Fname	Lname	Cno
"John"	"Smith"	"CS112"
"Mary"	"Smith"	"CS112"
"John"	"Brown"	"CS112"

Figure 16.3: Another example data for tables Professor and Teach.

The answer to such a modified query on the data of Figure 16.3 would again be empty. Yet, the answer to our question here should list Mary Brown.

Note that we seem to have run out of available syntax at this point- there does not seem to be any other reasonable query modification. Does that mean that we cannot ask such a question? The answer is: Of course, we can- we just need to recall and appropriately use one of the previously

introduced language features.

We must first determine what it is that we actually need. The concept of a composite key offers a clue: We do not want first names and last names to be treated separately. Rather, we want NOT IN to be tested between Fname/Lname pairs. In other words, we would like to be able to write a query as shown in Figure 16.4.

```
SELECT Fname, Lname
FROM Professor
WHERE      NOT IN
(Fname, Lname
  (SELECT Fname,
         Lname FROM
           Teach
        WHERE      "CS112"))))
(Cno =
```

Figure 16.4: Desired query form for the question Who does not teach CS112?

11

Negation Over Composite Keys

Unfortunately, this is syntactically incorrect: the IN and NOT IN operators cannot be used with a list of "pairs." They can, however, be used with a list of expression results!

In other words, even though the IN and NOT IN operators require that the subquery retrieve a single column in its SELECT clause, the specification for that column need not be a single basic attribute, but can be an expression.

Likewise, NOT IN can also be used with a list of expression results. All that is required is that the data types of these expressions be compatible.

The query of Figure 16.5 incorporates expressions to pose our question Who does not teach CS112?

```

SELECT Fname, Lname

FROM   or
Profess
WHERE   + Lname FROM
(Fname + Lname
NOT IN
(SELECT Fname Teach
      WHERE   "CS112")))
(Cno =

```

Figure 16.5: Using string concatenation expressions for the question Who does not teach CS112?

We assume here that columns Fname and Lname are implemented as fixed character strings, and that the concatenation operator + does not suppress extra spaces, so that Fname and Lname do not "intrude" into each other.

Using expressions in this manner, we can now pose questions involving "at most," "exactly," "only" and "either-or" (similar to questions E8 through E11) about professors teaching courses, as well as a question involving "youngest" (or "oldest") (similar to question E12) about

professors. These are left as exercises for the reader. **11**

THE ESSENCE OF SQL: A Guide to Learning Most
of SQL In the Least Amount of Time

17 Standard Questions Involving Every

Given an ability to use expressions around the IN and NOT IN operators, we can now pose question E13- Who takes every course?- as shown in Figure 17.1.

```

SELECT Sno
FROM Student
WHERE (SELECT
(Sno NOT Sno
IN
      FROM Student, Course
      WHERE (SELECT Sno
      (Sno + Cno + Cno FROM
      NOT IN      Take))))

```

Figure 17.1: Query Q13 for question E13 – Who takes every course?

This query, which is motivated by the rephrasing of question Q13 as Which students are not among those for whom there is a course that they do not take? (see Figure 9.1) is quite complex, and really has to be traced in order to understand how it works.

Because all column references in this query are local, this query is of Type 2. (The inner Sno and Cno come from the inner table Take, the middle Sno and Cno come from the middle tables Student and Course respectively, and the outer Sno comes from the outer table Student.) Therefore, it is executed inside-out.

We first evaluate the inner-most subquery, which returns a list of Sno/Cno concatenations from table Take.

.....  11

We then evaluate the middle subquery. This subquery executes a cross-product between tables Student and Course. Its NOT IN operator then retains only those wide rows from this cross-product where the Sno/Cno combinations are not retrieved by the inner subquery- i.e., not present in table Take- i.e., where student Sno does not take course Cno. The answer to the middle subquery then contains the Sno

values from these combinations- i.e., the student numbers of those students for whom there is a course that they do not take.

The condition of the outer query then fails for these students, leaving for the final answer the student numbers of only those students who are not among those retrieved by the middle subquery. In other words, those students who take every course .

.... - THE ESSENCE OF SQL: A Guide to
Learning Most of SQL In the Least Amount of Time

18 A Brief Interlude

We have just completed all of the standard questions from Figure 4.1, covering Type 1 and Type 2 queries, as well as SQL enhancement features. What remains are: the aggregation facilities, NULLs (remember the NULL and NOT NULL qualifiers from Figure 3.1?), the Type 3 queries, the DML

component of SQL- namely, the facilities to insert, delete and update rows, and SQL extensions.

11 THE ESSENCE OF SQL: A Guide to Learning
Most of SQL In the Least Amount of Time

19 Additional Standard Questions

Five additional standard questions, motivating SQL aggregation facilities, are listed in Figure 19.1.

E14. For each department that has more than 3 profes

sors older than 50, what is the average salary of such professors?

(A related question: For each department/rank combination that has more than 3 professors older than 50, what is the average salary of such professors?)

E15. What is the grade point average (GPA) of each student?

E16. What is the overall average salary of all professors who are older than 50?

E17. Whose (i.e., which professors') salary is greater than the overall average salary?

E18. Whose salary is greater than the average salary within that professor's department?

Figure 19.1: A list of additional standard questions motivating SQL aggregation facilities.

While the standard nature of these questions is not immediately apparent, it will become clear as we consider each question in turn.

11 Additional Standard Questions

11 THE ESSENCE OF SQL: A Guide to Learning
Most of SQL in the Least Amount of Time

20 Computing Aggregates for Groups

The query posing question E14- For each department that has more than 3 professors older than 50, what is the

average salary of such professors? - is shown in Figure 20.1.

```
SELECT Dept, AVG(Salary)
FROM Professor
WHERE      50)
  (Age >
GROUP BY Dept
HAVING (COUNT(.,.) > 3)
```


Figure 20.1: Query Q14 for question E14 - For each department that has more than 3 professors older than 50, what is the average salary of such professors?

This query involves two new clauses: the GROUP BY clause, which takes a list of columns as an argument, and the HAVING clause, which takes a condition as an argument. It also involves two aggregate functions: AVG() and COUNT(.,.).

Syntactically, the HAVING clause is positioned immediately after the GROUP BY clause. The two are placed after the WHERE clause, but before the ORDER BY clause (not present in this case).

Operationally, the following sequence of events takes place. First, SQL formulates and evaluates the query

```
SELECT ""
FROM Professor
WHERE (Age > 50)
```



This query is called the underlying query. It is formed from

our original query by dropping the GROUP BY and HAVING clauses and by replacing the SELECT specification with symbol "*" (asterisk). Thus, it retrieves all columns from all of the tables involved. (While, in this case, the underlying query uses only one table and is of Type 1, in general it can involve multiple tables and be of any type.)

The result of this underlying query is then grouped by the Dept column (or, more precisely, by the values remaining in this column). In other words, the rows are re-arranged (or, re-ordered) in such a way that all rows sharing the same Dept value are listed next to each other.

We note that grouping is not the same as sorting; in particular, grouping does not guarantee that among themselves the groups would be listed in any particular order. Thus, for example, a group of "Mathematics" rows may appear after the group of "English" rows, but before the group of "History" ones.

SQL then evaluates the HAVING clause condition ($\text{COUNT}(\text{" "}) > 3$) for each group, and eliminates those groups for which this condition fails. When evaluated for a group, the aggregate function $\text{COUNT}(\text{" "})$ returns the number of rows in that group. Thus; for our query, only those groups that have more than 3 rows in them remain.

Finally, SQL formulates one resulting row corresponding to each remaining group and retrieves it. In this case, according to the SELECT clause of our query

`SELECT Dept, AVG(Salary)`

this resulting row would contain the Dept value of the group, and the average of all its Salary values.

1. Formulate and evaluate the appropriate underlying query

```
SELECT ,  
FROM ...  
WHERE  
...
```

2. Group (re-arrange) the rows in the result of Step 1
according to the GROUP BY clause.

3. Evaluate the HAVING clause condition for each group, and
eliminate those groups for which this condition fails. 4. Formulate
one resulting row for each remaining group, according to the
SELECT clause, and retrieve it.

Figure 20.2: The evaluation mechanism for queries involving
aggregation.

We emphasize that this is a formal evaluation mechanism. As
such it may differ from, and in particular appear less efficient
than, the actual evaluation strategies taken by real SQL systems.
However, we do guarantee that it will always generate the same
result.

In addition to COUNT(,) and AVG(), SQL provides four other
aggregate functions: MIN(), MAX(), SUM(), and
COUNT(DISTINCT <column>), that also can be used in the
SELECT and HAVING clauses. The HAVING clause can also
involve Boolean operators NOT, AND and OR, as well as scalar
expressions. No subqueries are allowed in the HAVING clause,
however! These capabilities are illustrated by the query of Figure
20.3



```
SELECT Dept,  
        MIN(Salary), MAX(Salary), AVG(Salary),  
        COUNT(DISTINCT Salary)  
FR.()M Professor
```

```

WHERE      50)
  (Age >
GROUP BY Dept
HAVING (COUNT(*) > 3)
       OR (SUM(Salary) - 200000 > 0)

```

Figure 20.3: A query showing the use of all aggregate functions, and also showing the use of Boolean expressions in the HAVING clause.

For numeric arguments, the behavior of MIN(), MAX(), AVG() and SUM() is quite natural: given a column of values, they return the smallest value, the largest value, the average of the values, and the arithmetic sum of the values, respectively. (We note that the aggregate function AVG() does not suppress duplicate values in its argument.)

For non-numeric data types, functions AVG() and SUM() are generally not defined, and the behavior of MIN() and MAX() is consistent with the behavior of the less-than and greater-than operators for those data types. For example, function MIN() applied to a column of string values, would return the lexicographically smallest of them. (Recall that the sorting behavior imposed by the ORDER BY clause is also consistent with the behavior of the less-than and greater-than operators.)

Functions COUNT(*) and COUNT(DISTINCT <column>) are both counters, and apply to all data types. What they count, however, is quite different. As we have already noted, function COUNT(*) simply counts the number of rows in a group.

11

THE ESSENCE OF SQL: A Guide to

Learning Most of SQL In the Least Amount of Time

Function COUNT(DISTINCT <column>) counts how many different <column>-values are present in the group. For example, given a group with 5 rows containing Salary values (20000, 30000, 20000, 20000, 30000), function COUNT(DISTINCT Salary) would return 2- for the two

distinct values: 20000 and 30000. (While this use of DIS
TINCT serves a different purpose than its use in SELECT
DISTINCT, both uses are quite consistent with each other.)

The question corresponding to the query of Figure 20.3 can be phrased as follows: For each department that has more than 3 professors older than 50 or where the total salary of such professors exceeds \$200,000, what is the minimum salary, the maximum salary, the average salary, and the number of distinct salary values, of such professors?

SQL permits the omission of the HAVING clause if the condition on the groups is always True- i.e., all groups are desired for the answer. (Recall the similar ability with the WHERE clause.) This allows the question For each department, what is the average salary of those professors who are older than 50? to be posed as shown in Figure 20.4.

```
SELECT Dept, AVG(Salary)
FROM Professor
WHERE      50)
  (Age >
GROUP BY Dept
```

Figure 20.4: The query for the question For each department, what is the average salary of those professors who are older than 50?

If the argument to the GROUP BY clause is a list of several columns, then the grouping occurs for all of them simulta-

11 • Computing Aggregates for Groups

neously. To illustrate this, consider the query of Figure 20.5, which poses the question For each department/rank combination that has more than 3 professors older than 50, what is the average salary of such professors? (the related question to E1 4).

```
SELECT Dept, Rank, AVG(Salary)
```

```
FROM
```

```
    Professor
```

```
WHERE      50)
```

```
    (Age >
```

```
GROUP BY Dept, Rank
```

```
HAVING (COUNT(*) > 3)
```

Figure 20.5: SQL query for the question For each department/
rank combination that has more than 3 professors
older than 50, what is the average salary of such
professors?

In this query, rows are grouped both by Dept and by Rank in other words, all rows in a group must now agree both on Dept and on Rank values. (Generally, this creates more groups, each with fewer rows.) Since grouping does not imply sorting, the order of columns in the GROUP BY clause has no significance.

We note that in this case, the condition "... that has more than 3 professors older than 50 ..." will be applied to each department/ rank combination separately. Thus, it is possible for some department to survive this condition in query Q14, but to not survive this condition here, when the department is further subdivided by rank. (Find such a data example, and trace the two queries.)

Standard SQL imposes an extremely important syntactic restriction on the form of the SELECT and HAVING clauses in queries involving aggregate functions and/or the GROUP BY clause:
